

# Milestone 1: Data Pipeline

---

## QM 2023 Capstone Project

**Due:** Friday, Week 5 (February 20, 2026) by 11:59 PM **Points:** 50 (25% of capstone grade) **Format:** Team submission via GitHub Classroom

**Important:** Work in your team's **shared private GitHub repository** (the same repo you'll use for all capstone milestones). Do not create separate repositories for different milestones.

---

## Overview

Build the **foundational data infrastructure** for your capstone: integrate your primary dataset (REIT returns by default, or approved alternative) with supplementary data (economic indicators, policy measures, or market factors) into a clean, analysis-ready panel.

**Real-world context:** Professional analysts spend 60–80% of their time on data engineering. A robust pipeline ensures all subsequent analysis runs smoothly and reproducibly.

**Success criterion:** Another analyst (or your instructor) should be able to run your script from scratch and obtain identical results.

---

## Directory Structure

M1 outputs go to `data/final/`, not `results/`. Use `config_paths.py` for consistent paths:

- **Input:** Primary data from `data/raw/`; supplementary data as needed
- **Output:** Analysis-ready panel to `data/final/[dataset]_analysis_panel.csv` and `data/final/[dataset]_metadata.json`

See `STRUCTURE.md` in the Capstone Project folder for the full folder guide.

---

## Learning Objectives

1. **Fetch or load data** from appropriate sources (APIs, databases, CSV files)
  2. **Clean messy datasets** (missing values, duplicates, outliers, date formatting)
  3. **Merge datasets** on common keys (Date/Time, Entity ID) while preserving data integrity
  4. **Reshape data** into tidy panel structure (Entity × Time)
  5. **Document data decisions** (missing values, outliers, filters, transformations)
  6. **Produce reproducible output** (relative paths, clear comments, metadata)
- 

## Deliverables

1. Python Script: `capstone_data_pipeline.py`

### Requirements:

- Runs without errors from top to bottom
- Uses **relative paths only** (no hardcoded `C:\Users\....`)
- Includes clear section headers and comments explaining each step

- Outputs a tidy CSV file ready for analysis

### Core Sections:

```
"""
QM 2023 Capstone Project: Milestone 1 Data Pipeline
Team: [Your Team Name]
Members: [List names]
Date: [Submission date]
Dataset: [REIT Master / Cryptocurrency / Housing / Other]

This script fetches, cleans, and merges [your dataset] into a tidy panel structure.
"""

# Section 1: Imports and setup
# Section 2: Load primary dataset
# Section 3: Clean data (missing values, outliers, duplicates)
# Section 4: Fetch/load supplementary data (economic indicators, policy measures)
# Section 5: Merge primary + supplementary data
# Section 6: Reshape to panel structure (Entity × Time)
# Section 7: Save tidy output and metadata
```

## 2. Tidy Output File: [\[dataset\\_name\]\\_analysis\\_panel.csv](#)

### Generic Panel Structure:

[entity_id]	[time_var]	[outcome]	[char_1]	[char_2]	[driver_1]	[driver_2]	[control_1]	[control_2]
[ID_1]	[T_1]	[Y_1]	[X_1]	[X_2]	[Z_1]	[Z_2]	[C_1]	[C_2]

**REIT:** permno | ym | ret | mcap | sector | price | fedfunds | mortgage30us | cpiaucsl | unrate

**Crypto:** token\_id | date | return\_pct | mcap\_usd | token\_type | volume\_usd | reg\_severity | btc\_return | cpi | vix

**Housing:** fips\_code | year\_quarter | rent\_growth\_yoy | median\_price | region | income\_median | mortgage\_rate | gdp\_growth | hpi | unrate

### Requirements:

- **Long format** (one row per entity-time observation)
- **No missing keys** (entity\_id and time\_var must be non-null for all rows)
- **Date consistency** (time\_var: **YYYY-MM** monthly, **YYYY-MM-DD** daily, **YYYY-QX** quarterly)
- **Merged correctly** (supplementary data aligned to time, no duplicate rows)

## 3. Data Quality Report: [M1\\_data\\_quality\\_report.md](#)

### Required sections:

1. **Data Sources** — Primary: name, source, coverage (N entities, frequency, date range), initial row count, key variables. Supplementary: source, variables, date range.
2. **Data Cleaning Decisions** — For each: variable, % missing/count, decision (drop/impute/cap/winsorize), justification. Cover: missing values, outliers, size/volume filters, duplicates, data type corrections.
3. **Merge Strategy** — Join type (left/inner), keys, alignment. Before/after row counts; verification that counts make sense.

4. **Final Dataset Summary** — Entity variable, time variable, balanced/unbalanced, final dimensions. Sample statistics table (mean, std, min, max, missing %). Data quality flags.
5. **Reproducibility Checklist** — Script runs, relative paths, output location, no manual editing, metadata + AI Audit complete.
6. **Ethical Considerations** — What data are we losing? Who might we exclude? Example: "By dropping mcap < \$10M REITs, we exclude micro-caps. Acceptable for institutional focus; small REITs have different risk profiles." Transparency; test alternatives in M3 robustness.

**Example entries (adapt to your dataset):**

Decision	REIT	Crypto
Missing outcome	Drop (5.2%—delistings)	Drop (2.1%—new listings)
Outliers	Winsorize 99th/1st pct	Cap ±100%
Size filter	Drop mcap < \$10M	Drop tokens <30 days history
Duplicates	Keep first	Average price, sum volume

**Worked REIT example (Section 1–4):**

- **Primary:** REIT Master Panel, instructor-provided. 532 REITs, monthly 2015-01–2023-12, 56,800 rows. Keys: permno, ym, ret, mcap, sector, price.
- **Supplementary:** FRED (FEDFUNDS, MORTGAGE30US, CPIAUCSL, UNRATE) via pandas-datareader, 108 months.
- **Cleaning:** ret 5.2% missing → drop (delistings); sector 3.1% → drop; returns >200% or <-100% → winsorize 99/1 pct; mcap <\$10M → drop; duplicate permno-ym → keep first; ym "2020m1" → "2020-01".
- **Merge:** Left join on month; 48,258 rows before and after; 0 NaNs in supplementary.
- **Final:** 532 REITs × 108 months = 48,258 obs (unbalanced).
- **Sample stats table:** Variable | Mean | Std Dev | Min | Max | Missing (%). Fill for outcome, chars, drivers, controls.

**Sign-off:** All team member names.

---

4. AI Audit Appendix: [AI\\_AUDIT\\_APPENDIX.md](#)

**Required for all milestones.** Document AI use with "Disclose, Verify, Critique":

```

## AI Tools Used
- [ ] ChatGPT / Copilot / Claude / Other

## Per Task
- **Task:** [Description]
- **Prompt:** "[Exact prompt]"
- **AI Output:** [Code/text provided]
- **Verification:** [How you tested]
- **Critique:** [Right/wrong, your corrections]

## Summary
Total AI use, primary use cases, verification method.
Responsibility: All code is tested and our responsibility.

```

**No AI Audit Appendix = 0/50** (enforced strictly). See M4 memo template for detailed examples.

---

# Technical Requirements

## 1. Data Sources

### REIT Teams (Default):

- **Primary:** `data/reit_master_raw.csv` — permno, ym, ret, mcap, sector, price. Known issues: ~5% missing returns, ~2% duplicate permno-month, outliers, inconsistent ym formatting.
- **Factors:** `data/reit_factors.csv` — ym, SIZE, VALUE, MOM, QLTY, LOWVOL, REV, ewtret, vwtret (clean).
- **Supplementary:** FRED via `pandas-datareader` — FEDFUNDS, MORTGAGE30US, CPIAUCSL, UNRATE (match REIT date range).

**Alternative Dataset Teams:** Principles: Entity × Time structure; supplementary drivers/controls; merge on time frequency; document all cleaning decisions.

---

## 2. Data Cleaning Pipeline (Generic Pattern)

### Step 1 — Load & inspect:

```
primary_df = pd.read_csv(RAW_DATA_DIR / '[your_dataset].csv')
primary_df.columns = primary_df.columns.str.lower().str.replace(' ', '_')
print(primary_df.shape, primary_df.head(), primary_df.info())
```

### Step 2 — Parse time:

```
# Monthly: primary_df['ym'] = pd.to_datetime(primary_df['ym'].str.replace('m', '-'),
errors='coerce')
# Daily:    primary_df['date'] = pd.to_datetime(primary_df['timestamp'], unit='s')
```

**Step 3 — Missing values:** Document before (`df.isnull().sum()`). Typically drop outcome; justify imputation.

```
print(f"Missing outcome: {primary_df['outcome'].isnull().sum()}")
({100*primary_df['outcome'].isnull().mean():.1f}%)")
primary_clean = primary_df.dropna(subset=['outcome'])
```

**Step 4 — Duplicates:** `drop_duplicates(subset=['entity_id', 'time_var'], keep='first')` (or average/sum per domain).

**Step 5 — Outliers:** Winsorize, cap, or drop. Justify by domain.

```
upper, lower = primary_clean['outcome'].quantile(0.99),
primary_clean['outcome'].quantile(0.01)
primary_clean['outcome'] = primary_clean['outcome'].clip(lower=lower, upper=upper)
# Crypto/high-vol: primary_clean['outcome'] = primary_clean['outcome'].clip(-100, 100)
```

**Step 6 — Size/volume filters:** Drop entities below threshold (e.g., mcap < \$10M). Document.

---

### 3. Fetch Supplementary Data

#### FRED (REIT, Housing, Macro):

```
import pandas_datareader as pdr
federal_funds = pdr.DataReader('FEDFUNDS', 'fred', start, end)
mortgage = pdr.DataReader('MORTGAGE30US', 'fred', start, end)
# ... concat, reset_index, rename columns. Align time to match primary (e.g.,
# .dt.to_period('M')).
```

**Events/Custom:** Load CSV, ensure datetime, fill missing dates (e.g., severity=0 for no event).

### 4. Merge Strategy

Align supplementary time to primary. Use **left join**. Fill missing: forward fill, 0 for "no event," or drop.

```
# Align time: econ_df['ym'] = pd.to_datetime(econ_df['DATE']).dt.to_period('M')
merged = primary_clean.merge(econ_df, on='ym', how='left')

# Critical check
assert merged.shape[0] == primary_clean.shape[0], "Row count mismatch!"
# If merged >> primary: many-to-many. supp_df.drop_duplicates(subset=['time_var']) before
merge
```

#### Panel verification:

```
print(merged['entity_id'].nunique(), merged['time_var'].nunique(), merged.shape[0])
obs_per_entity = merged.groupby('entity_id')['time_var'].count()
print(obs_per_entity.min(), obs_per_entity.max()) # unbalanced if min != max
```

### 5. Save Output

```
merged.to_csv(FINAL_DATA_DIR / '[dataset]_analysis_panel.csv', index=False)
metadata = {
    'dataset': '[name]', 'n_entities': int(merged['entity_id'].nunique()),
    'n_time_periods': int(merged['time_var'].nunique()), 'n_obs': int(merged.shape[0]),
    'time_range': f'{merged["time_var"].min()} to {merged["time_var"].max()}',
    'variables': list(merged.columns),
    'cleaning_decisions': {'missing_outcome': 'Dropped', 'outliers': 'Winsorized'},
    'size_filter': '...'}
with open(FINAL_DATA_DIR / '[dataset]_metadata.json', 'w') as f:
    json.dump(metadata, f, indent=2, default=str)
```

### Grading Rubric (50 points)

Component	Points	Criteria
<b>Reproducibility</b>	15	Script runs; relative paths; clear section headers and comments
<b>Data Cleaning</b>	12	Missing, duplicates, outliers handled; before/after counts documented
<b>Merge Integrity</b>	8	No row loss/duplication; supplementary data aligned correctly
<b>Panel Structure</b>	8	Correct entity-time format; verified dimensions
<b>Documentation</b>	7	Data quality report complete; metadata JSON; cleaning decisions justified

#### Zero-Credit Conditions:

- Missing `AI_AUDIT_APPENDIX.md` = **0/50**
- Script won't run (syntax/path errors) = **0/50**
- Hardcoded absolute paths = **-10 points**

## Common Pitfalls

Pitfall	Problem	Solution
Hardcoded paths	<code>C:/Users/...</code> won't run on other machines	Use <code>config_paths</code> and relative paths
Merge duplicates	merged rows 10x primary	One row per time in supplementary; deduplicate before merge
Time type mismatch	string vs datetime; NaT after merge	<code>pd.to_datetime()</code> on both datasets
No cleaning docs	"Cleaned data" with no details	Document every step: counts + economic justification
Missing AI Audit	No appendix	<b>0/50</b> — log AI use as you work
Wrong thresholds	REIT winsorize applied to crypto	Use domain-appropriate thresholds for your data

## Testing Checklist (Run Before Submission)

- Script runs from scratch (delete `data/final/`, rerun)
- No absolute paths (search for `C:\, /Users/`)
- No manual Excel editing — script alone produces output
- Row counts: merged ≤ primary (no accidental duplication)
- Spot-check supplementary values (e.g., FEDFUNDS Jan 2020 ~1.55%)
- All deliverables: `capstone_data_pipeline.py`, `*_analysis_panel.csv`, `*_metadata.json`, `M1_data_quality_report.md`, `AI_AUDIT_APPENDIX.md`
- Team names in all files
- Alternative dataset approved (if applicable) by Week 4

## Submission Instructions

1. Accept M1 assignment via GitHub Classroom
2. Clone repo, add files, commit, push:

```
git add capstone_data_pipeline.py data/final/*.csv data/final/*.json  
M1_data_quality_report.md AI_AUDIT_APPENDIX.md  
git commit -m "M1: Data Pipeline - [Dataset Name]"  
git push origin main
```

3. Verify files visible on GitHub; GitHub Actions will run checks.

**Deadline:** Friday, Week 5 (Feb 20, 2026) 11:59 PM | **Late:** 10% per day, up to 3 days.

---

## Resources and Support

- **Office Hours:** Dr. Seagraves, Mon & Wed 3–5 PM (Helm 122-D)
- **Starter script:** [starter/capstone\\_data\\_pipeline.py](#)
- **pandas-datareader:** [Official docs](#)

**Debugging:** Merge not working? Print both DataFrames; ensure keys match (type + values). API failing? FRED can rate-limit; try `time.sleep(1)` between requests.

---

## Next Steps After M1

M2 (Week 9) — EDA and visualization. **Don't wait until Week 8 to start.** Common M2 discovery → M1 fix:

- M2 shows "too many missing in lag 2" → Add 2-period lag in M1
  - M2 sector analysis fails → Verify sector cleaning in M1
  - M2 correlations look strange → Check merge/date alignment in M1
- 

QM 2023: *Statistics II, Spring 2026, University of Tulsa | Dr. Cayman Seagraves (cayman-seagraves@utulsa.edu)*